

COVID-19 Twitter Dataset with Latent Topics, Sentiments and Emotions Attributes

Important notes

1. The dataset described in this paper is available for download at OpenICPSR: <https://doi.org/10.3886/E120321>
2. The dataset license is also available at the OpenICPSR download folder. Essentially, the dataset license considers the need to be consistent with Twitter's terms as the dataset is built upon content provided by Twitter API, and [CC BY-NC 2.0](#).
3. Hence, if you intend to use the dataset, you should read and agree with Twitter's [Terms of Service](#), [Privacy Policy](#), [Developer Agreement](#), and [Developer Policy](#). The user should also read the [restricted uses](#) from Twitter to avoid using the dataset for any potentially inappropriate use.
4. In compliance with Twitter's content redistribution terms, our released dataset only includes "tweet_ID" and "user_ID" from Twitter's original data attributes. Users shall be able to use "tweet_ID" and/or "user_ID" to retrieve or "hydrate" the other attributes (such as the actual "text", "tweet_created_at", "retweet_count", "location", "followers_count") through the Twitter API directly.
5. No redistribution beyond the user's immediate research group or lab is allowed. The user shall direct any interest or request to the corresponding author.
6. Citation of this paper:
Gupta, R., Vishwanath, A., and Yang, Y. (2020), COVID-19 Twitter Dataset with Latent Topics, Sentiments and Emotions Attributes, *Preprint* at: <https://arxiv.org/abs/2007.06954>
7. Enquiries can be sent to the corresponding author via [email](#).

COVID-19 Twitter Dataset with Latent Topics, Sentiments and Emotions Attributes

Raj Kumar Gupta, Ajay Vishwanath, Yinping Yang

Affective and Social Intelligence Group, Social and Cognitive Computing, Institute of High Performance Computing, Agency for Science, Technology and Research (A*STAR), Singapore
Corresponding author: Yinping Yang (yangyp@ihpc.a-star.edu.sg)

Abstract

This paper presents a large annotated dataset on public expressions related to the COVID-19 pandemic. Through Twitter's standard search application programming interface, we retrieved over 63 million coronavirus-related public posts from more than 13 million unique users since 28 January to 1 July 2020. Using natural language processing techniques and machine learning based algorithms, we annotated each public tweet with seventeen latent semantic attributes, including: 1) ten binary attributes indicating the tweet's relevance or irrelevance to ten detected topics, 2) five quantitative attributes indicating the degree of intensity of the valence or sentiment (from extremely negative to extremely positive), and the degree of intensity of fear, of anger, of sadness and of joy emotions (from extremely low intensity to extremely high intensity), and 3) two qualitative attributes indicating the sentiment category and the dominant emotion category, respectively. We report basic descriptive statistics around the topics, sentiments and emotions attributes and their temporal distributions, and discuss its possible usage in communication, psychology, public health, economics and epidemiology research.

Background & Summary

The 2019 coronavirus disease (COVID-19) was first officially reported as acute respiratory infections caused by an unknown virus in Wuhan city, Hubei province in China on 31 December 2019. To date, the disease has infected 10,357,662 people worldwide and claimed 508,055 lives, according to the World Health Organization (WHO)'s situation report as at 1 July 2020 [1].

The issues surrounding the pandemic are increasingly challenging and complex. The complexity not only comes from the disease itself, but also by the surge of the medical, social, behavioural and economic issues that the disease has brought about, such as reports on daily counts of new cases and mortality rates, scientific discoveries, government responses, news reporting of social behaviours such as panic buying and food hoarding, impact on businesses and economic outlook, and changes in people's everyday lives. The challenges are multi-faceted and unprecedented. There is a growing recognition of the need for multidisciplinary research efforts to support the COVID-19 pandemic response, including disciplines such as social and behavioural science [2] and mental health science [3].

Twitter is a popular microblogging site widely used by Internet users. According to Statista, as of the fourth quarter of 2019, Twitter had 152 million active users worldwide [4]. Twitter provides the research community a rich source of information about when, where and what people have to say in their posts (known as "tweets") through its free, publicly accessible standard application programming interface (API) service. However, the raw tweet content is mostly in textual format and is not readily analysable. When there is a huge amount of tweets, it takes a significant amount of time to accurately extract information related to people's concerns, feelings and emotions for human analysts and researchers to process and analyse for in-depth patterns and insights.

Despite the fact that Twitter has opened up real-time, full-fidelity data streams related to COVID-19 tweets since late April 2020 [5], and that a few recent studies have leveraged Twitter for COVID-19 studies (e.g., [6, 7]), to the best of our knowledge, no others have provided a tweet-by-tweet research resource with rich, semantically and psychologically meaningful attributes surrounding the topics, sentiments and emotions from the tweets content. The dataset described in this paper provides tweet-by-tweet tagging of the topic clusters that a tweet is semantically related to, the sentiment it is expressing, and the emotional properties it is associated with.

Figure 1 presents the data schematic structure and the data processing methods.

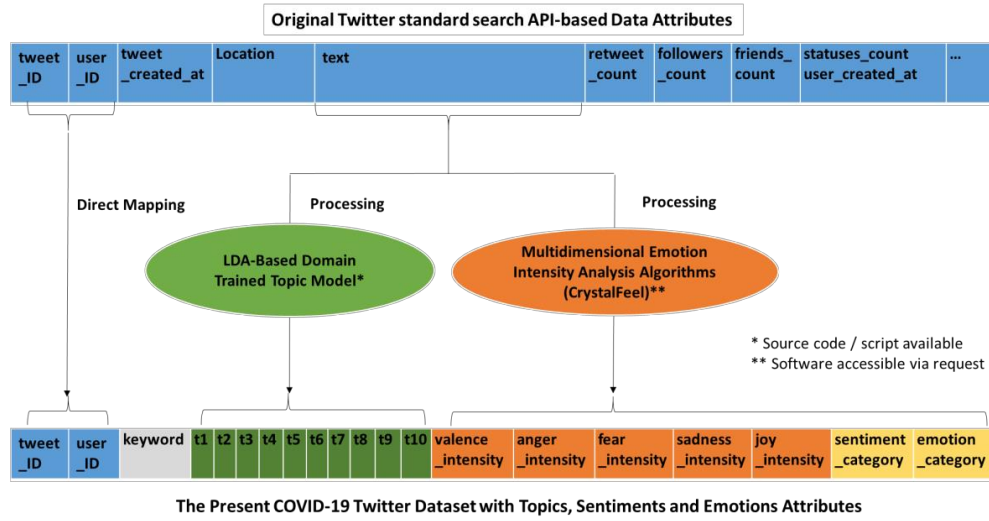


Figure 1. Overview of the data structure and processing methods

An initial analysis of the world’s emotion trends using a part of this dataset has shed light on the significant change of people’s emotional responses to the pandemic from late January to early April [8]. This initial study focused on the “emotion_category” attribute solely, and found that anger has overtaken fear as the dominant emotion in tweets.

We seek to make the full dataset with tweet-by-tweet topics, sentiments and emotions attributes available for broad research communities. The purpose is to allow researchers to perform more in-depth investigations in all possible areas such as discovering the correlational patterns between with other important variables such as government policy communications, demographics, economic indicators, as well as epidemiological markers.

Methods

We describe the methods behind the data collection and data processing with a focused interest on tracking and understanding the latent topics, sentiments and emotions surrounding the COVID-19 pandemic. We applied natural language processing (NLP) techniques, in particular, statistical topic clustering techniques that detect tweets surrounding similar topic clusters. We also applied previously machine learning-trained and validated algorithms to tag each tweet with a sentiment valence (unpleasantness / pleasantness), as well as intensity scores of four different emotions – anger, fear, sadness and joy.

Collection of raw Twitter data related to COVID-19

We started to set up our data collection app early February 2020 by querying Twitter's standard search API [9] using three keywords "corona", "wuhan" (many people refer to the virus as "wuhan virus" at initial stages when official names were not yet announced), and "nCov" (WHO first named the virus as "2019-nCoV"). On 11 February 2020, upon WHO officially renamed the disease as "COVID-19", we added "covid" as a new search keyword. Simple sharing of these tweets (i.e., retweets) are not collected for the dataset.

The Twitter API returns the tweet text content with a rich range of useful attributes. For example, we were able to download the following 12 attributes in our local database.

- **tweet_ID**: The unique identifier for this tweet
- **tweet_created_at**: The UTC time when this tweet was created
- **retweet_count**: Number of times this tweet has been retweeted
- **text**: The actual UTF-8 text of the tweet
- **favorite_count**: Number of times this tweet has been liked
- **hashtag_text**: Name of the hashtag, minus the leading '#' character
- **user_ID**: The unique identifier for the user who created the tweet
- **followers_count**: Number of followers this account currently has
- **friends_count**: Number of users this account is following
- **statuses_count**: Number of tweets (including retweets) issued by the user
- **user_created_at**: The UTC datetime that the user account was created on Twitter
- **location**: The user-defined location in this account's profile

Processing for the topic attributes

Although all the retrieved tweets are relevant to at least one of the four COVID-19 related keywords, there are many facets or subtopics people have covered in the tweets' "text" content. We applied an unsupervised topic clustering technique called Latent Dirichlet Allocation (LDA) to facilitate the understanding of the subtopics. LDA is a probabilistic generative model which learns a multinomial distributions of latent topics in a given document and words in a given topic [10]. The advantage of LDA is that it is independent of the corpus size which makes it algorithmically efficient to learn topic clusters within a corpus with a large number of tweets such as ours.

First, we pre-processed each raw tweet by converting it to ASCII characters, removing accented characters, forming bigrams and trigrams, filtering out stop words (including most rare and most frequent words), and performing text tokenization. These pre-processed tweets were then converted into a bag-of-words (BoW) corpus. The training data's date range is 28 January to 27 May 2020, which consists of 51 million tweets.

Next, we randomly sampled 1% of the BoW corpus and trained an LDA model whose inference was performed using online variational Bayes [11]. Using the trained LDA-based topic model, we obtained 100 topic clusters. The following list illustrates the top ten topics detected in the dataset (e.g., "topic 1") and the ten most representative words associated with each detected topic (e.g., "people, cases, new, deaths, time, china, realdonaldtrump, lockdown, trump" for "topic 1"), respectively.

- **topic 1**: people, cases, new, deaths, time, china, realdonaldtrump, lockdown, trump
- **topic 2**: health, help, people, need, think, vaccine, care, fight, support
- **topic 3**: pandemic, f**k, months, killed, question, wait, looks, trump, impact
- **topic 4**: pay, donate, lie, focus, song, gates, page, google, caused
- **topic 5**: florida, drink, named, nature, marketing, pr, ncdcgov, farmers, cr

- **topic 6:** rules, bed, drtedros, speaks, privacy, parliament, physicians, strength, joke
- **topic 7:** dies, pmoindia, ndtv, ai, narendramodi, mohfwindia, shoot, drharshvardhan, battle
- **topic 8:** ye, ke, behaviour, brought, hidden, yup, smell, zero hedge, odds
- **topic 9:** excuse, humanity, salary, wind, gtgt, rats, ice, beard, mosque
- **topic 10:** internet, allah, teacher, dance, el, rona, weed, crush, fk

Lastly, for each tweet in the entire dataset, we assigned a relevance label (“1” or “0”) using the trained LDA model based on the contribution of each topic (over a total of 100 topic clusters) to the tweet (“1” indicates if the contribution is > 1%, where “0” indicates otherwise).

Table 1 shows example tweets that are tagged with the corresponding topic clusters that they are relevant to, respectively. The first ten examples show tweets that are solely relevant to each of the ten topic clusters. The last example shows that a tweet can be relevant to multiple topic clusters simultaneously.

Table 1. Examples of the tweets and their corresponding topic attributes

Example tweet text	t1	t2	t3	t4	t5	t6	t7	t8	t9	t10
Remember when doja cat said corona was just the flu	1	0	0	0	0	0	0	0	0	0
STAY SAFE from CORONAVIRUS There is currently no antiviral treatment or vaccine to prevent Coronavirus (COVID-19)	0	1	0	0	0	0	0	0	0	0
No. It is a part of One Belt One Road. #ChinaVirus WHO says Wuhan coronavirus outbreak is not yet a pandemic	0	0	1	0	0	0	0	0	0	0
Bill Gates: This is how long it may take before Americans can be completely safe from COVID-19	0	0	0	1	0	0	0	0	0	0
Florida’s COVID-19 website guru blasts bosses, hints at data suppression	0	0	0	0	1	0	0	0	0	0
Municipality closes 336 shops in all six governorates for violating corona rules	0	0	0	0	0	1	0	0	0	0
@Naveen_Odisha @narendramodi @MoHFW_INDIA @PrakashJavdekar @drharshvardhan @dr_arunsahoo I request to postpone our...	0	0	0	0	0	0	1	0	0	0
Previously its DAM funds. now ita corona fund. pochna sirf ye tha k DAM fund kaha tk pohancha	0	0	0	0	0	0	0	1	0	0
apparently child abuse cases are going up and the excuse is that parents are stressed over covid-19. im sorry but	0	0	0	0	0	0	0	0	1	0
today i found out weed kills the corona virus and i ain’t been worried since	0	0	0	0	0	0	0	0	0	1
PICS: Built in 10 days – 1000 bed hospital opens to battle coronavirus in Wuhan All 4 Women	1	1	1	0	1	0	0	0	0	0

Processing for the sentiment and emotion intensity attributes

As sentiments and emotions are subjective information embedded in the unstructured “text” content, it is a necessity to extract such information with targeted tools. We used CrystalFeel [12], a collection of five machine-learning based algorithms to extract the sentiment and emotions scores. The development of CrystalFeel involved training and experimental evaluations of features derived from affective lexicons, parts-of-speech, and word embeddings [13], using tweets manually annotated with intensity ground truth values [14].

Table 2 shows five example tweets tagged with these five attributes. The first example shows a tweet with a moderate (i.e., neither very negative nor very positive) sentiment “valence_intensity”, and the other four examples show tweets with relatively high intensity score for joy, anger, fear and sadness dimensions. It shall be noted that, in certain cases such as the fourth example (“Being higher risk of covid has me all over the place. Appt is in about 90 mins. Im scared, worried and anxious”), the intensity score may be exceeding the 0-1 normal ranges, which indicates that these cases represented extreme intensities beyond the algorithms’ original training samples.

Table 2. Examples tweets and their corresponding sentiment valence and emotions intensity scores attributes

Example tweet text	valence_ intensity	anger_ intensity	fear_ intensity	sadness_ intensity	joy_ intensity
Community hospital Bright Vision transfers all patients to make room for stable COVID-19 cases	0.505	0.391	0.444	0.423	0.334
Seeing this face for just a few minutes on her 97th birthday today made my heart so happy! #Greatful #COVID #Nana	0.930	0.143	0.177	0.209	0.822
To any fellow nationalist celebrating boris getting the corona virus are nothing less than an absolute cunt.	0.318	0.665	0.520	0.495	0.260
Being higher risk of covid has me all over the place. Appt is in about 90 mins . Im scared, worried and anxious	0.231	0.542	1.074	0.670	0.174
When u cant handle shit any longer and u feel hopeless #COVID_19 starting to make my depression worse	0.076	0.617	0.836	0.940	0.077

Processing or converting intensity to the sentiment category attribute

To facilitate more straightforward interpretation, we also converted CrystalFeel algorithm’s sentiment valence intensity score (quantitative value as “valence_intensity”) into to a categorical label (qualitative value as “sentiment_category”) using the following logic.

```
// # Initialize the sentiment category in a "neutral or mixed" class
1     sentiment_category = "neutral or mixed";
// # Assign the sentiment category based on the degree of the valence intensity
2     if(valence_intensity <= 0.30):
3         sentiment_category = "very negative";
4     elif(valence_intensity < 0.48):
5         sentiment_category = "negative";
6     elif(valence_intensity > 0.70):
7         sentiment_category = "very positive";
8     elif(valence_intensity > 0.52):
9         sentiment_category = "positive";
```

Table 3 shows the five tweets examples tagged with their corresponding sentiment categories, qualitatively indicating the sentiment each tweet is mainly expressing.

Table 3. Examples of the tweets data and their corresponding sentiment categories

Example tweet text	sentiment_category
Community hospital Bright Vision transfers all patients to make room for stable COVID-19 cases	neutral or mixed
To any fellow nationalist celebrating boris getting the corona virus are nothing less than an absolute cunt.	negative
Being higher risk of covid has me all over the place. Appt is in about 90 mins . Im scared , worried and anxious	very negative
When u cant handle shit any longer and u feel hopeless #COVID_19 starting to make my depression worse	very negative
Seeing this face for just a few minutes on her 97th birthday today made my heart so happy! #Greatful #COVID #Nana	very positive

Processing or converting intensity to the emotion category attribute

As the underlying dominant emotion behind the sentiments carries more information than the overall valence or sentiment, we also converted CrystalFeel's joy, anger, fear and sadness intensity scores into an "emotion_category" value to facilitate interpretation.

To achieve this, we applied a logic that leverages all the valence and emotions intensities scores from CrystalFeel's outputs. We first use "valence_intensity" as the first-line criterion as this dimension that has very high accuracy, i.e., 0.816 in terms of Pearson correlation with human annotated ground truth values. We then use the relative intensity comparing the three primary negative emotions, anger, fear and sadness to assign a corresponding dominant emotion category. The following script describes the conversion logic:

```
// # Initialize the sentiment category in a "no specific emotion" class
1     emotion_category = "no specific emotion";
// # Assign the emotion category when valence intensity score exceeds 0.52
2     if(valence_intensity > 0.52):
3         emotion_category = "joy or other positive expression";
// # Assign the emotion category when valence intensity score falls below 0.48
4     elif(valence_intensity < 0.48):
5         emotion_category = "anger";
6     if((fear_intensity > anger_intensity) and (fear_intensity >= sadness_intensity )):
```

```

7      emotion_category = "fear";
8      elif((sadness_intensity > anger_intensity) and sadness_intensity > fear_intensity)):
9      emotion_category = "sadness";

```

Table 4 shows the five tweets examples tagged with the dominant emotion categories.

Table 4. Examples of the tweets data and their corresponding emotion category results

Example tweet text	emotion_category
Community hospital Bright Vision transfers all patients to make room for stable COVID-19 cases	no specific emotion
To any fellow nationalist celebrating boris getting the corona virus are nothing less than an absolute cunt.	anger
Being higher risk of covid has me all over the place. Appt is in about 90 mins . Im scared , worried and anxious	fear
When u cant handle shit any longer and u feel hopeless #COVID_19 starting to make my depression worse	sadness
Seeing this face for just a few minutes on her 97th birthday today made my heart so happy! #Greatful #COVID #Nana	joy or other positive expression

It is useful to note that the above-mentioned conversion logic turning CrystalFeel’s original emotional intensity scores to “sentiment_category” and “emotion_category” labels are based on application assumptions where CrystalFeel is used for processing *short informal text* (e.g., tweets, Facebook posts and comments). The conversion thresholds are derived from heuristics and social media corpora we continuously monitor in our research (see more in [12]).

Users may define and adjust their own conversion logic as far as it is appropriate or suitable for different applications. For example, for converting the emotional intensity scores to meaningful categories on *short formal text* (e.g., news headlines), the conversion logic shall be adjusted accordingly.

Data Records

As at 1 July 2020, our system collected a total of 63,347,100 English tweets worldwide using the four COVID-19 related keywords, with the first retrievable date being 28 January 2020.

The data record is constructed to include two comma-separated value (CSV) files:

1. tweets_topics_sentiments_emotions (5k sample).csv

This is a very small, randomly selected sample of the full dataset. It can be used to quickly understand the data structure and attributes using any typical applications such as Microsoft Excel.

2. tweets_topics_sentiments_emotions.csv

This is the full 63 million tweets csv file. The file is approximately 8 GB. It is recommended to use python+pandas to view and retrieve data record in this file.

Here, each of the data record file has a total of 20 columns (i.e., attributes), including: the “tweet_ID” (the unique identifier for this tweet), the “user_ID” (the unique identifier for the user), the “keyword” (“corona”, “wuhan”, “nCov” or “covid”, which we used to query the

Twitter API to obtain the corresponding tweet), and the following 17 latent topics, sentiments and emotions related attributes.

- **t1:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t2:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t3:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t4:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t5:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t6:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t7:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t8:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t9:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **t10:** A binary value of 0 or 1, where 0 – this tweet is not relevant to this topic; 1 – this tweet is relevant to this topic
- **valence_intensity:** A real-valued score with its normal range as from 0 to 1, where 0 - this text expresses extremely high intensity of unpleasant feelings; 1 - this text expresses extremely high intensity of pleasant feelings
- **anger_intensity:** A real-valued score with its normal range as from 0 to 1, where 0 - this text expresses extremely low intensity of anger; 1 - this text expresses extremely high intensity of anger
- **fear_intensity:** A real-valued score with its normal range as from 0 to 1, where 0 - this text expresses extremely low intensity of fear; 1 - this text expresses extremely high intensity of sadness
- **sadness_intensity:** A real-valued score with its normal range as from 0 to 1, where 0 - this text expresses extremely low intensity of sadness; 1 - this text expresses extremely high intensity of sadness
- **joy_intensity:** A real-valued score with its normal range as from 0 to 1, where 0 - this text expresses extremely low intensity of joy; 1 - this text expresses extremely high intensity of joy
- **sentiment_category:** A categorical value indicating the sentiment a tweet is expressing, with a value being either “very negative”, “negative”, neutral or mixed”, “positive”, “very positive”
- **emotion_category:** A categorical value indicating the most dominant emotion a tweet is expressing, with a value being either “anger”, “fear”, “sadness”, “joy or other positive expression”, “no specific emotion”

Important note. In compliance with Twitter’s content redistribution terms, the original Twitter data attributes in our released dataset only include “tweet_ID” and “user_ID”. Users shall be able to use “tweet_ID” and/or “user_ID” to retrieve or “hydrate” the other attributes (such as the actual “text”, “tweet_created_at”, “retweet_count”, “location”, “followers_count”) through the standard search API from Twitter directly [9].

Technical Validation

Raw tweets coverage

As at 1 July, our system has collected a total of 63,347,100 English tweets worldwide using the four COVID-19 related keywords, with the first retrievable date being 28 January 2020.

Table 6 presents the data overview. Out of all the tweets, majority are retrieved based on the “covid” keyword which returned 44,930,122 tweets, or 71%. On average, 17,029 COVID-related tweets were posted every hour, or 408,691 tweets every day.

Table 6. Overall Twitter data overview (28 January – 1 July 2020; 155 days)

Keyword used	Volume / no. of tweets collected	Distribution
corona	15,492,526	24%
wuhan	2,569,674	4%
ncov	354,778	1%
covid	44,930,122	71%
total	63,347,100	100%

These tweets are posted by 13,749,346 unique users based on the “user_ID” attribute.

Nevertheless, it is useful to note that a known limitation of Twitter’s standard search API is that it does not guarantee the retrieved tweets are exhaustive, due to indexing and other reasons. In other words, the search API retrieves *relevant* but *not all* the tweets that match the search keywords.

Geographical representativeness

To assess the geographic representativeness of the tweets, we converted the “location” from the original Twitter attribute into a “country_region” attribute. This is done using GeoNames’ cities15000 geographic database [15], which contains a mapping between all cities with a population > 15000 or capitals and a country code. For example, original location “Ontario, Canada” is converted to country_region code as “Canada”, “India” is converted to as “India”, “Shanghai” is converted to as “Kuwait”, “London” is converted to as “United Kingdom”. If the location is indicated as “online”, “The Entire Universe!” or left blank (i.e., no match can be found using the GeoNames database), the country_region is coded as “-”.

The geographical coverage of the tweets is estimated to contain users coming from more than 170 countries, regions and territories worldwide. Figure 2 shows a visualization of the geographical representativeness.

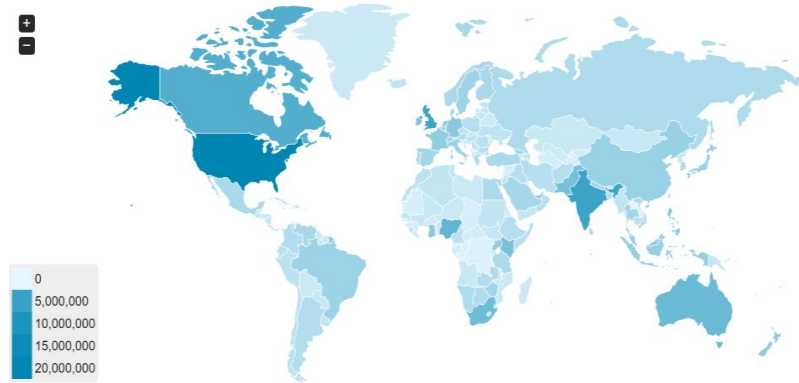


Figure 2. Geographical coverage and volume distribution of the tweets

Validity of the processing methods

Topic identification. The quality of the topic model was evaluated using metrics including perplexity and coherence scores based on suggestions from the literature [16]. We obtained the top ten topics, i.e., “t1”, “t2”, ..., “t10”, that received a relatively high coherence scores (c_v measure, mean = 0.575) from a model optimized by learning 100 topics and hyper-parameters α as a fixed normalized asymmetric Dirichlet prior ($1/\text{topic_number}$) and $\eta = 0.909$. We obtained ten topics out of 100 extracted from 500,000-odd data points, which is a 1% sample from the full dataset. Conceivably, training an LDA-based topic model with data from specific twitter accounts, smaller and more focused date range, particular countries of interest, or certain hashtags, would yield more targeted and meaningful results. Hence, we provide our Python source code to help researchers quickly apply and adapt the model for further usage scenarios.

Emotional intensity scores. The accuracy in determining “valence_intensity”, “fear_intensity”, “anger_intensity”, “sadness_intensity”, and “joy_intensity” are systematically validated in prior research [13], and are subsequently tested for predictive validities in other NLP tasks [17, 18, 19]. As reported in Gupta and Yang’s original evaluation experiments using out-of-training-sample test data, the CrystalFeel algorithms’ accuracies in terms of Pearson correlation coefficient (r) with manually annotated test data, are 0.816 on valence intensity, and are 0.708, 0.740, 0.700 and 0.720 on joy intensity, anger intensity, fear intensity and sadness intensity [13]. The predictive validity of the valence, joy, anger, fear and sadness intensity scores on other tasks has been studied and demonstrated in the context of predicting news social popularity in Facebook and Twitter [17], in predicting the ingredients of happy moments [18], and in detecting propaganda techniques in news articles [19]. Hence, researchers may examine the use of the sentiment and emotions intensity scores directly without conversions.

Sentiment category and emotion category labels. The “sentiment_category” and “emotion_category” attributes are obtained based on a conversion logic or “codebook” presented in the “Methods” section. The conversion principle that allows each tweet to be labelled with one of the five emotion categories (i.e., “fear”, “anger”, “sadness”, “joy or other positive expression”, “no specific emotion”) follows a conceptual simplification that a single dominant emotion exists for each tweet, though some tweets may express “mixed emotions” [20], such as express anger and fear simultaneously. Other conversion logic may be explored in future research. For example, Mohammad et al. [14] suggest to use the mid-scale threshold, i.e., 0.5, to differentiate high-intensity vs. non-high-intensity emotions. Researchers shall examine the intended applications and determine the conversion threshold accordingly.

Topics coverage

We checked the tweets volume related to the top ten identified topic clusters. A vast majority of tweets were related to two or more topics, which form 60% of the total tweets. The tweets that *solely* pertained to “t1” have the highest volume, consisting of 14,467,272 tweets or 23% out of the total data volume.

Table 7 presents the overall tweet topics statistics. Figure 3 depicts with a visualization of the topic clusters in the context of the tweets volume.

Table 7. Overall tweets distribution over topic clusters

Tweets related to attributes	Volume / no. of tweets	Distribution
two or more topics	38,065,393	60%
t1 only	14,467,272	23%
t2 only	7,345,398	12%
t3 only	2,328,557	4%
t4 only	703,455	1%
t5 only	94,969	0.15%
t6 only	64,236	0.10%
t7 only	56,718	0.09%
t8 only	35,003	0.06%
t9 only	8,890	0.01%
t10 only	31,075	0.05%
other single topics	146,134	0.23%
total	63,347,100	100%

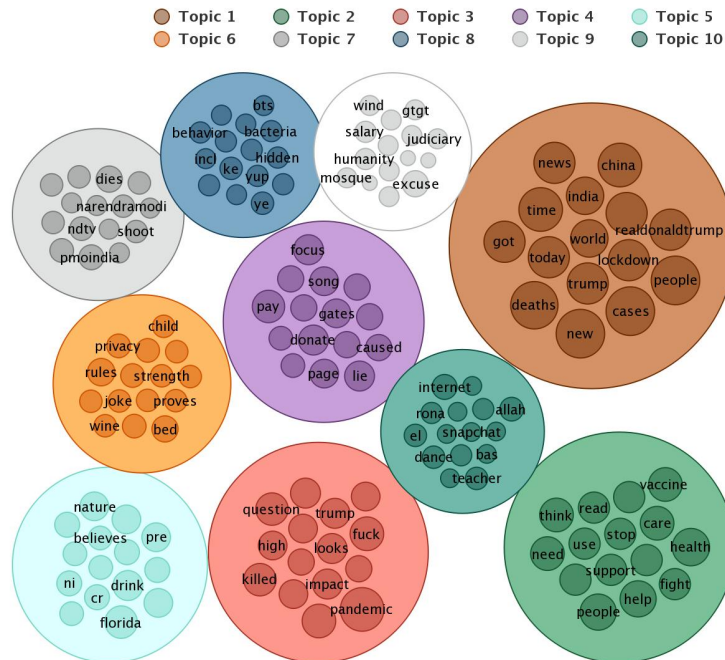


Figure 3. The ten topic clusters showing different issues surfaced in the total tweets

Sentiment intensity and sentiment category coverage

The quantitative “sentiment_intensity” averaged for the full dataset is 0.460, with the most negative tweet having its valence intensity score of -0.058, and the most positive tweet having its valence intensity score of 0.989 (Table 8).

Table 8. Sentiment intensity descriptive statistics

Attribute name	Mean	Standard deviation	Median	Min	Max
valence_intensity	0.460	0.093	0.461	-0.058	0.989

Qualitatively, the counts and distributions for valence intensity score converted into sentiment categories counts are presented in Figure 4. The results indicate that the majority of the tweets are “negative” or “very negative”, forming 58% of the total 63,347,100 tweets.




Figure 4. Five sentiment categories tweets count and distribution

Plotting the “sentiment_category” values over *daily* aggregated tweets counts suggested more nuanced patterns (see Figure 5). For example, the single-day peak during this period was 1,075,087 tweets (629,938 were “negative” tweets and 65,165 were “very negative” tweets), which took place on 13 March 2020, one day immediately following WHO’s announcement on the disease as a “pandemic”. Further analysis may look into, for example, the sentiment changes before and after more targeted time periods based on critical announcements (e.g., to study a week before and after 13 March 2020). The dataset may also allow for further research to explore the correlations and predictive values based on the sentiment and emotion scores, when over-laid with economic indicators (e.g., stock market changes).




Figure 5. Five sentiment categories in stacked chart showing their evolvement over time

Emotion intensity and dominant emotion category coverage

Using the four quantitative emotions intensities attributes, overall statistics show that “fear_intensity” has the highest mean value of 0.440, closely followed by “anger_intensity” of 0.437. Table 9 reports the descriptive statistics for the four emotions intensity scores.

Table 9. Emotion intensity descriptive statistics

Attribute	Mean	Standard deviation	Median	Min	Max
fear_intensity	0.440	0.096	0.438	-0.053	1.118
anger_intensity	0.437	0.085	0.433	0.028	1.100
sadness_intensity	0.411	0.078	0.405	0.039	0.968
joy_intensity	0.305	0.091	0.302	-0.091	0.936

Qualitatively, the counts and distribution of the various most dominant emotion categories based “emotion_category” attribute are presented in Figure 6. The results suggest that, over the five months in total, tweets that are dominantly expressing “anger” (16,864,419 tweets, 27%) and tweets that are dominantly expressing “fear” (16,668,156 tweets, 26%) formed the majority of the total tweets.




Figure 6. Count of tweets and distribution of each dominant emotion category

We checked the daily counts of the four emotions for the five-month period (see Figure 7). The significance of the change can be illustrated using the contrast of results at the start and at the end of our data range. For example, as at 28 January, a total of 23,405 tweets were posted for the day, and the tweets with “anger” as the dominant emotion formed 15% of the total 23,405 tweets, far less than those tweets with “fear” as the most dominant emotion which formed 53% of the total 23,405 tweets. In contrast, as at 1 July, a total of 432,736 total tweets were posted for the day, and the tweets with “anger” as the most dominant emotion formed 30% of the 432,736 tweets, exceeding those tweets with “fear” as the most dominant emotion, which formed 23% of the 432,736 tweets). The trends surfaced some interesting patterns: While both “fear” and “anger” dominated in the overall counts, the trend plot shows that over time, the relative distribution of “fear” has been decreasing and the relative distribution of “anger” has been increasing; meanwhile, “joy and other positive expressions” have been increasing, though in a slower rate (See [8] which provides an interpretation based on analysis of early time coverage of this dataset).




Figure 7. Daily distribution of the four primary emotions over time

Usage Notes

This paper presents a large COVID-19 tweets dataset with additional, psychologically meaningful attributes. This dataset may create opportunities to understand both global and local conversations and social sentiments in real time, at a large scale, potentially leading to very rich insights on human behaviours and behavioural changes surrounding the unprecedented pandemic. We envisage its potential usage in five broad areas.

First, for media and mass communication research, the dataset can be useful for communication scientists and professionals in evaluating and improving government response, policies and media communications towards the unprecedented pandemic crisis. For example, a recent study compared communications efforts of health authorities in the United States, the United Kingdom and Singapore on Facebook during the early period of COVID-19 [21]. As the virus continuously hit different countries in different timeframes and governments implemented different response strategies and policies, one direction of an ongoing related work is to overlay the location attribute, examine and compare sentiments, emotions and topics associated with different countries (e.g., [22]). The dataset may also help to study how media’s topical and emotional framing in their headlines and titles are different from those expressed by the general public.

Second, the dataset is of inherent interest for psychology research. The granularity of the tweets metadata in terms of “tweet_created_at” and “user_ID” may allow researchers to dive deeper of the more nuanced trends with deeper psychological accounts and insights. One possibility is to look into the sentiments and emotions differences over more fine-grained timelines, examine cultural differences, and segregate the users which are influencers vs the general public. Future research may also look into leveraging user characteristics inference techniques (e.g., [23]) and the present dataset to investigate user community-specific tendencies and issues.

Third, as the pandemic escalates in its severity and geographical span, and is likely to last for a prolonged period, public mental health issues (e.g., [24]) are more prevalent. The dataset may be used to examine public mental wellbeing. Prior literature (e.g., [25, 26]) has established the linkage between fear (as an emotion) and anxiety (as a mental disorder), and between sadness (as an emotion category) and depression (as a mental disorder). Hence it can be fruitful to study the value of the emotion intensity scores and their trends in the context of its duration, frequency and in relation to various user communities.

Fourth, it is potentially useful to overlay publicly available economic indicators (e.g., daily stock market data, monthly unemployment rates reports), and investigate in more depth on how the twitter topics, sentiments and emotions trends present predictive value in future research.

Last but not least, data scientists and epidemiology researchers may find the dataset useful. For example, prior research in Zika [27] and other infectious disease outbreaks [28] have studied and found useful insights in overlaying with air travel networks and virus genome. Hence the dataset may contribute towards revealing more hidden patterns and relationships of the large-scale social media content and other pandemic-related data streams.

Code Availability

The source scripts for the trained LDA-based topic model are available at our GitHub page: https://github.com/ajvish91/covid_twitter_scripts. Info on CrystalFeel can be accessed from: <https://socialanalyticsplus.net/crystalfeel>. Its access is available upon request from the corresponding author.

Acknowledgements

This work is supported by the Agency for Science, Technology and Research (A*STAR) under its A*ccelerate Gap Fund (ETPL/18-GAP050-R20A & ETPL/18-GAP052-R20A), the Singapore Ministry of Health's National Medical Research Council under its COVID-19 Research Fund (COVID19RF-005) and the Institute of High Performance Computing, A*STAR under its Core Fund. The authors thank Anita Sheldenkar, Brandon Loh, Cui Mengyang, Gangeshwar Krishnamurthy, Lim Keng Hui, Jiahui Lu, Joseph JP Simons, May Oo Lwin, Nur Atiqah Othman, Paul E. Cain, Sebastian Maurer-Stroh, Therese Quieta, Thomas Vuong, Tuan Le Mau, Wan Kum Seong, Wong Chi Kit, and Zhang Mila, for helpful discussions. We are grateful for the help from Nur Atiqah Othman for her proofreading which helped to enhance the clarity of the paper. All errors that remain are our sole responsibility.

Author Contributions

RG acquired the data and extracted sentiment and emotions features. AV extracted topic clusters features. YY initiated, conceptualized, and led the manuscript. All authors performed data analysis, contributed to the manuscript writing, reviewed the content and agreed with the submission.

Competing Interests

The authors declare the following competing interests: RG and YY are co-inventors of the CrystalFeel tool which was used to extract the sentiment and emotion related attributes. No other conditions or circumstances present a potential conflict of or competing interest for the other authors.

References

1. World Health Organization, *Coronavirus disease 2019 (COVID-19) Situation Report – 163 on 1 July 2020* https://www.who.int/docs/default-source/coronaviruse/situation-reports/20200701-covid-19-sitrep-163.pdf?sfvrsn=c202f05b_2 (2020).
2. Bavel JJV, Baicker K, Boggio PS, Capraro V, Cichocka A, Cikara M, et al. Using social and behavioural science to support COVID-19 pandemic response. *Nat Hum Behav.* **4**, 460–471(2020).
3. Holmes EA, O'Connor RC, Perry VH, Tracey I, Wessely S, Arseneault L, et al. Multidisciplinary research priorities for the COVID-19 pandemic: a call for action for mental health science. *Lancet Psychiatry* (2020).

4. Clement, J. *Number of monetizable daily active Twitter users (mDAU) worldwide from 1st quarter 2017 to 1st quarter 2020*
<https://www.statista.com/statistics/970920/monetizable-daily-active-twitter-users-worldwide> (2020).
5. Twitter, *COVID-19 Stream* <https://developer.twitter.com/en/docs/labs/covid19-stream/overview> (2020).
6. Abd-Alrazaq A, Alhuwail D, Househ M, Hamdi M & Shah Z. Top Concerns of Tweeters During the COVID-19 Pandemic: Infection Study. *J Med Internet Res.* **22** (2020).
7. Budhwani H & Sun R. Creating COVID-19 Stigma by Referencing the Novel Coronavirus as the “Chinese virus” on Twitter: Quantitative Analysis of Social Media Data. *J Med Internet Res.* **22** (2020).
8. Lwin MO, Lu J, Sheldenkar A, Schulz PJ, Shin W, Gupta R, & Yang Y. (2020). Global Sentiments Surrounding the COVID-19 Pandemic on Twitter: Analysis of Twitter Trends, *JMIR Public Health Surveill* **6** (2020).
9. Twitter *standard search API*
<https://developer.twitter.com/en/docs/tweets/search/overview/standard> (2020).
10. Blei, D. M., Ng, A. Y., & Jordan, M. I. Latent dirichlet allocation. *Journal of Machine Learning Research*, **3**, 993-1022 (2003).
11. Hoffman, M., Bach, F. R. & Blei, D. M. Online learning for latent dirichlet allocation. In *Advances in Neural Information Processing Systems*, 856-864 (2010).
12. CrystalFeel - Multidimensional Emotional Intensity Analysis. *Institute of High Performance Computing, A*STAR*. <https://socialanalyticsplus.net/crystalfeel> (2020).
13. Gupta, R.K. & Yang, Y. Understanding and Detecting Intensity of Emotions using Affective Lexicons, In *Proceedings of the 12th International Workshop on Semantic Evaluation (NAACL-SemEval)* (2018).
14. Mohammad, S., Felipe Bravo-Marquez, Mohammad Salameh & Svetlana Kiritchenko. SemEval-2018 task 1: Affect in tweets. In *Proceedings of the 12th International Workshop on Semantic Evaluation (NAACL-SemEval)*. (2018).
15. GeoNames, *GeoNames Geographical Database*
<http://download.geonames.org/export/dump/cities15000.zip> (Accessed on 15 Feb 2020)
16. Chang, J., Gerrish, S., Wang, C., Boyd-Graber, J. L. & Blei, D. M. Reading tea leaves: How humans interpret topic models. In *Advances in Neural Information Processing Systems*, 288-296 (2009).
17. Gupta, R.K. & Yang, Y. Predicting and Understanding News Social Popularity with Emotional Salience Features, In *Proceedings of the 27th ACM International Conference on Multimedia 2019 (ACM MM)* (2019).
18. Gupta, R.K., Bhattacharya, P. & Yang, Y. What constitutes happiness? Predicting and characterizing the ingredients of happiness using emotion intensity analysis, In *Proceedings of the AAAI-19 Workshop on Affective Content Analysis (AAAI-AFFCON)* (2019).
19. Krishnamurthy, G., Gupta, R.K. & Yang, Y. Detecting propaganda in news articles using sentence-level emotional salience features, In *Proceedings of the 14th International Workshop on Semantic Evaluation (COLING-SemEval)* (2020). Forthcoming. Preprint at: <https://arxiv.org/abs/2008.13012>
20. Berrios, R., Totterdell, P., & Kellett, S. Eliciting mixed emotions: a meta-analysis comparing models, types, and measures. *Frontiers in Psychology*, **6**, 428 (2015).
21. Sesagiri Raamkumar A, Tan SG & Wee HL. Measuring the Outreach Efforts of Public Health Authorities and the Public Response on Facebook During the COVID-19 Pandemic in Early 2020: Cross-Country Comparison. *J Med Internet Res* **22** (2020).
22. Lwin MO, Lu J, Sheldenkar A, Schulz PJ, Shin W, Gupta R & Yang Y. The evolution of public sentiments during the COVID-19 pandemic: A case comparison of Italy, Singapore, South

Korea, and the United States, Working Manuscript (under review at *Frontiers in Public Health*)

23. Volkova, S., Bachrach, Y., Armstrong, M. & Sharma, V., March. Inferring latent user properties from texts published in social media. In *Proceedings of Twenty-Ninth AAAI Conference on Artificial Intelligence* (2015).
24. Luo, M., Guo, L., Yu, M. & Wang, H. The Psychological and Mental Impact of Coronavirus Disease 2019 (COVID-19) on Medical Staff and General Public – A Systematic Review and Meta-analysis, *Psychiatry Research* **291** (2020).
25. Öhman, A., Fear and anxiety. *Handbook of Emotions*, 709-729. (2008).
26. Karp, D.A., *Speaking of sadness: Depression, disconnection, and the meanings of illness*. (Oxford University Press, 2017).
27. Grubaugh ND, Saraf S, Gangavarapu K, et al. Travel Surveillance and Genomics Uncover a Hidden Zika Outbreak during the Waning Epidemic, *Cell* **178** 1057-1071 (2019).
28. Tatem AJ, Rogers DJ & Hay SI. Global transport networks and infectious disease spread. *Adv Parasitol* **62** 293–343 (2006).